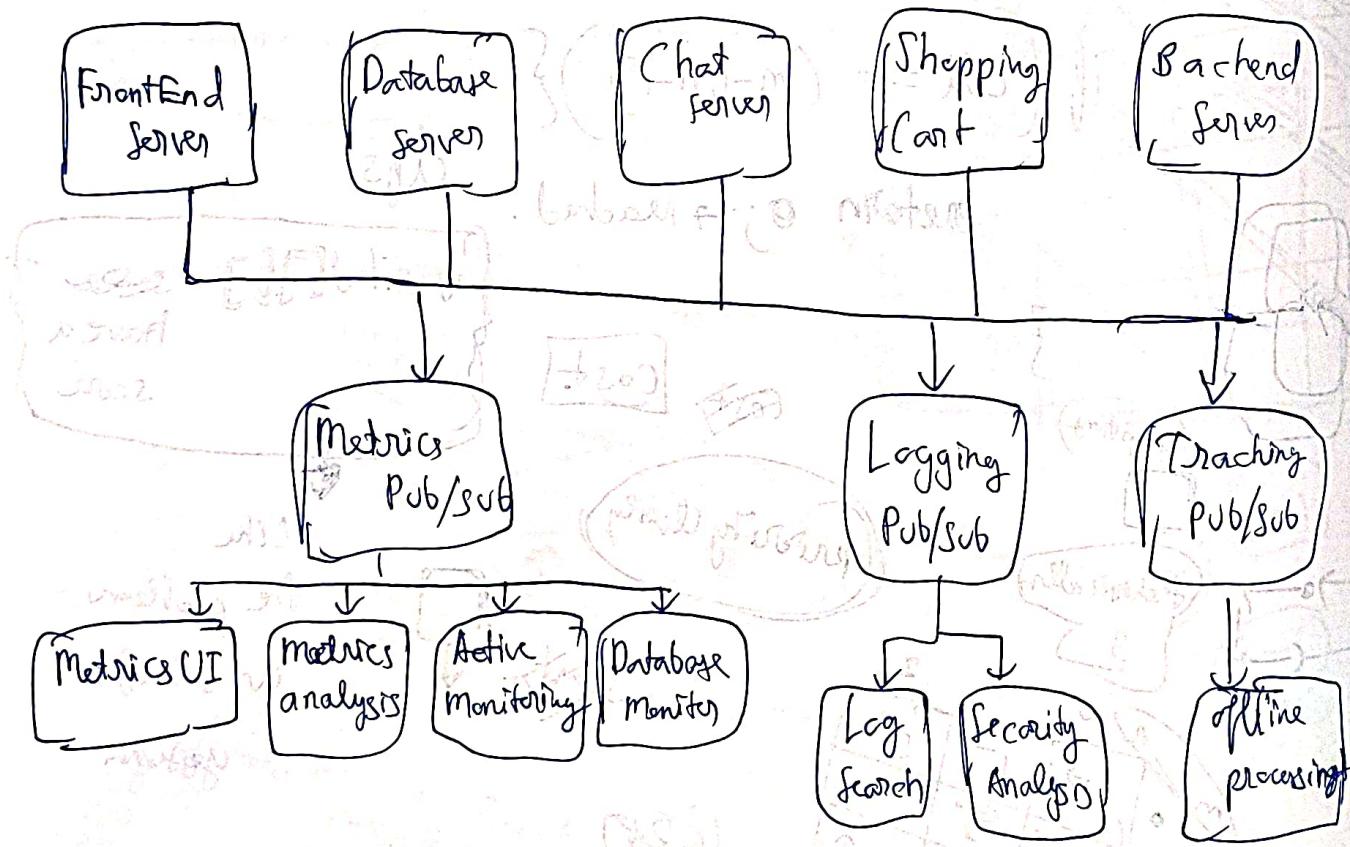
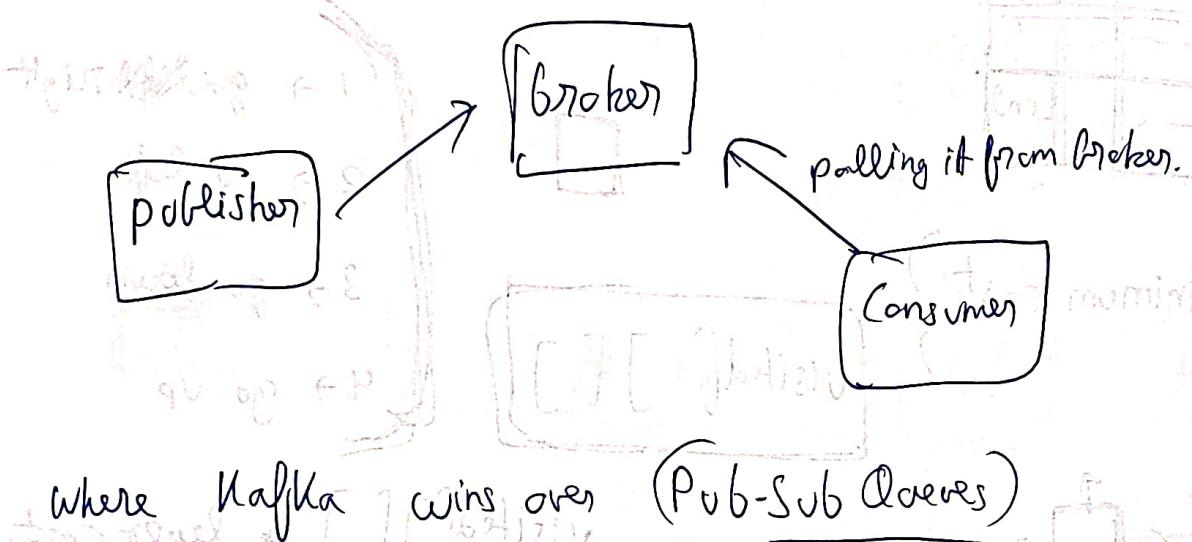
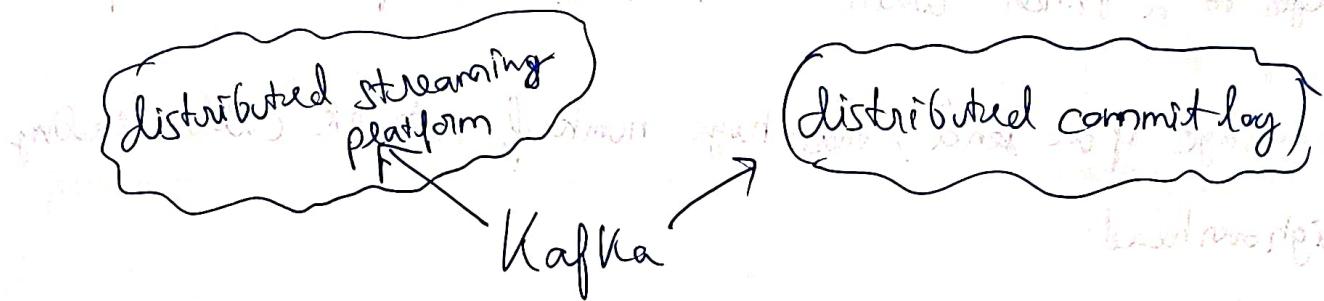


Chapter 1



Same duplicate data which has to be sent across to different queues for different consumption processes and requirements

Kafka solves this problem by becoming a centralised "distributed commit log".



Similar to a file system, or a database (commit log), Kafka makes this durable and data can be read in the way we want. + there is added replication to handle failures.

(Messages & Batches)

From a database background think of message as a row/record.

Message (row/record) is batched and written all in parallel at a time.

Message has an optional day key.

(Key, message)

Key is also a byte array

No use of key, unless key is used to generate which partition for message is destined for.

($\text{hash}(\text{key}) \% \text{numPartitions}$) → deciding where to write

Given that #Partitions are not changing, ensures messages with the same key are written to the same partition.

Batching is a process of putting multiple messages before sending them to a broker. Before sending messages to a broker waits till it sends the info to a topic.

Every message if we send, then huge number of network calls leading to high overhead.

So batch messages and send together to broker. We will learn more about these settings in chapter 2.

- ① Compression is applied on these batches to compress it and reduce payload size to send over network.

Schemas

In Kafka the messages are stored in (Array of byte). Chunks of Data. Kafka doesn't store message as

key = "K1"
value = {
 name = " -- "
 topic = " -- "

Kafka needs bytes, same as network, over network we don't need to send bytes and not the message as it is.

But at the producer and consumer side, we need to serializer and deserializer respectively to process the message accordingly.

JSON → Javascript Object Notation

XML → Extensible Markup Language

These can be used for schema definitions, but have some problems or rather are not compatible among versions.

[
JSON 1.0 (Pub) & JSON 2.0 (Sub)]

→ Problems with backward compatibility

Apache Avro is used for this Schema formatting. Avro is a serialization framework providing both backward & forward compatibility for both.

Topics & Partitions

In terms of database, Topics = {database table / file in a FileSystem}

Partition is same as how we partition a table in database.

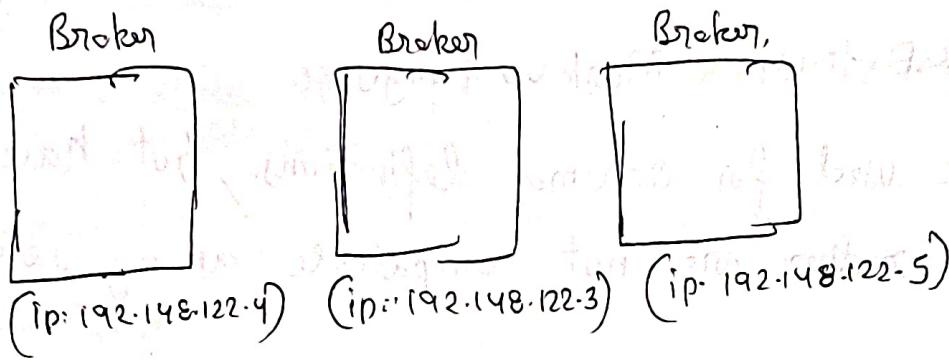
Kafka → "distributed commit log".

Messages are written to it in append only fashion.

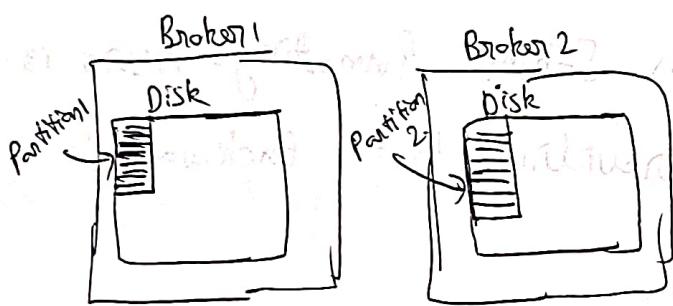
Ordering present within single partition,

Ordering not present within topic. (If multiple partitions).

Here I had a concept clearing enlightenment.



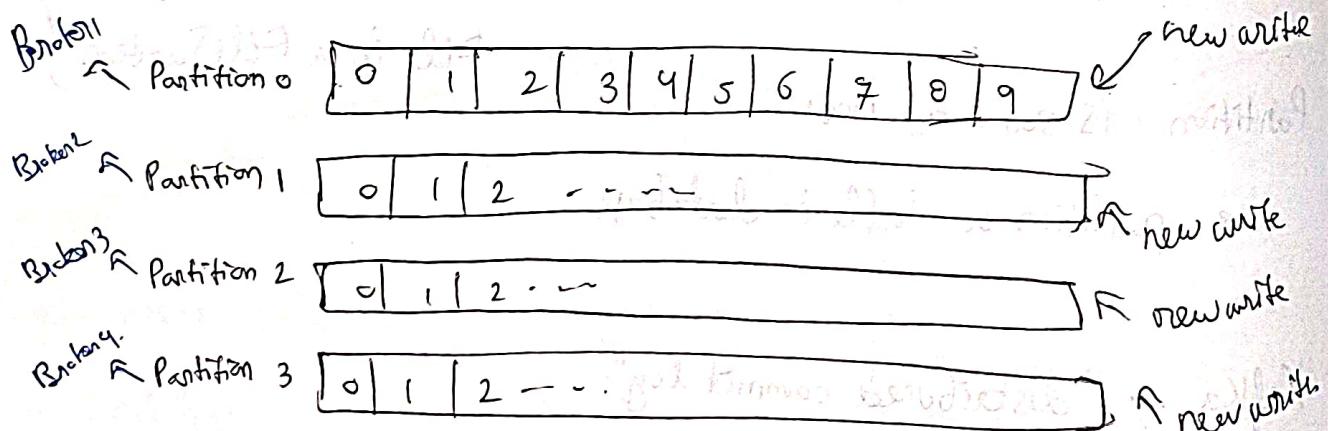
These are the actual servers in Kafka - Real machines with Ram and disk.



Inside each broker the File System is divided into multiple ~~partitions~~ segments, which act as partitions.

And group of partitions together is called a topic.

Topic Name = "topicName".



Each of these ~~topics~~ partitions can be hosted on different brokers / servers. (Actual servers), which can be individually scaled differently.

Scale up brokers as per load on partition.

(Producers & Consumers)

(Producer)

By default, when a message is produced to a topic, Kafka will try to balance messages across a partition evenly.

But we can write logic to direct message to a particular partition.

$(\text{Key}, \text{message}) \rightarrow \text{hash}(\text{Key}) \rightarrow \text{to direct message to a partition}$

(Consumer)

When producer produces a message to a partition, it gets automatically added to the Kafka.

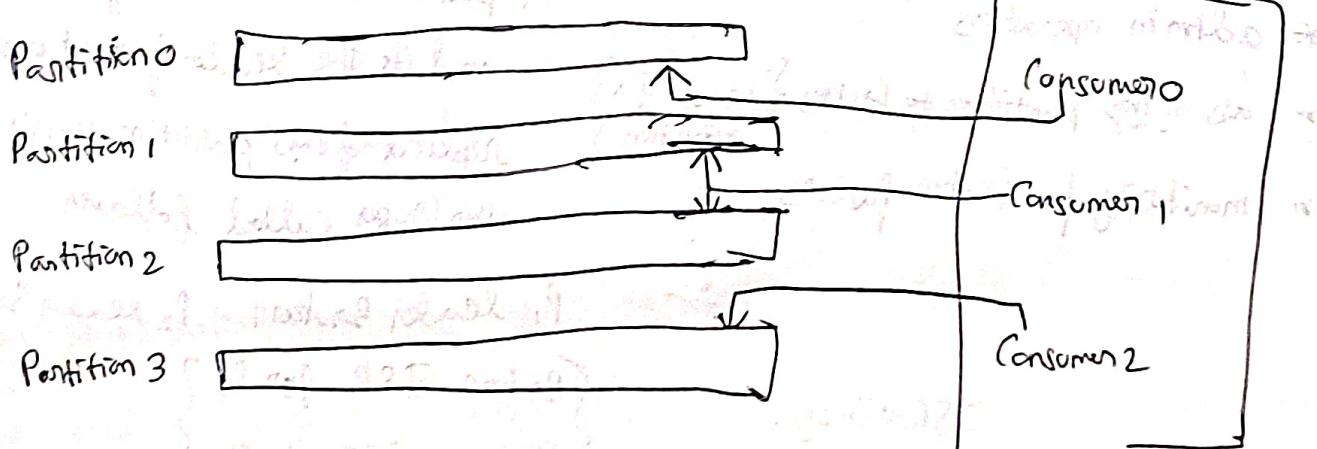
And Kafka adds an offset metadata \rightarrow for each message before adding it to the partition.

Each message within a partition is given an offset, unique and is greater than the previous offset, though its not monotonically increasing.

Consumer works as a part of consumer group.

One partition \rightarrow to one consumer only assigned

one consumer group



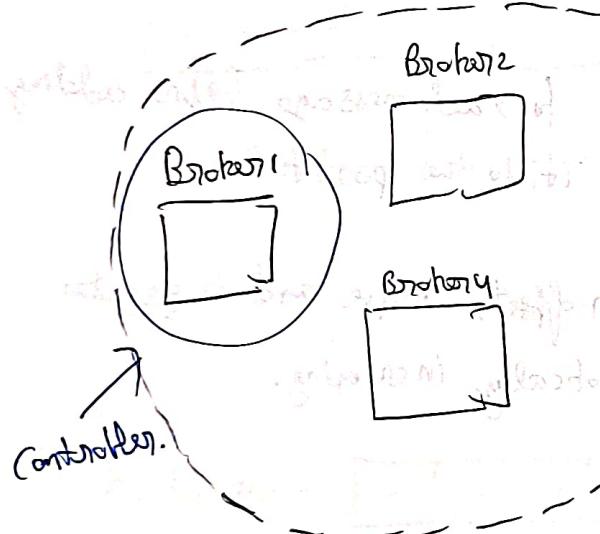
Brokers and Clusters

A single Kafka server is called a broker, the Broker receives messages from producers and assigns offsets to them and write messages to store on disk.

Broker \rightarrow [Single instance = Broker]
(ip)

Single Broker \rightarrow Scale, ^{on} vertical scaling can handle

[1K partitions & 1M msgs/s].



Cluster Controller does

- * admin operations
- * assigning partitions to brokers } Leader of a partition
- * monitoring for broker failures

ISR \rightarrow In sync
Replica

1 partition is owned by 1 broker and it's the leader of that partition, replica of this partition exists on other broker called follower.

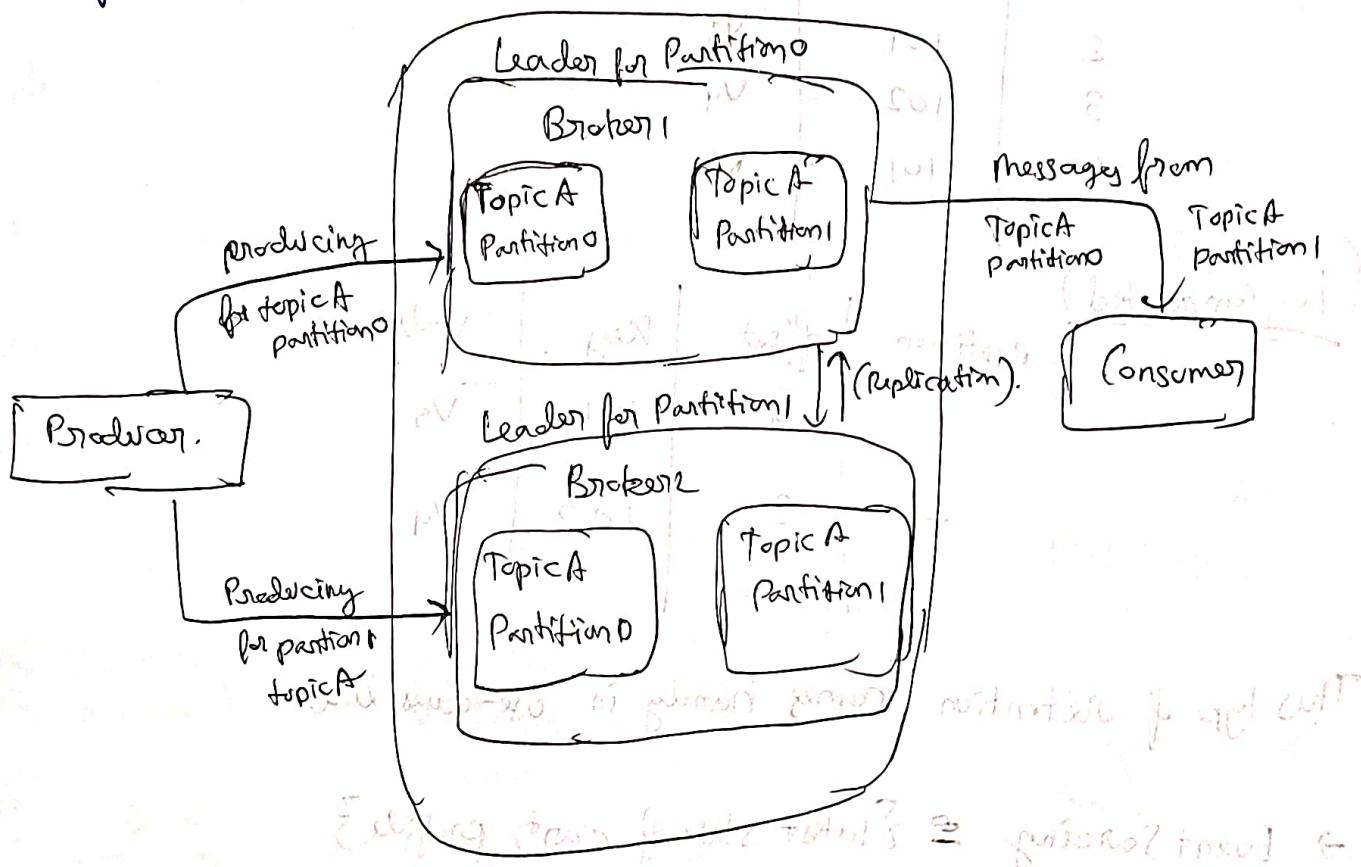
P₁ leader Broker 1 P₂ leader Broker 2
 { Broker 2 ISR for P₁ } { Broker 1 ISR for P₂ }

Producer produces messages to only leader broker of that partition.

Remember → Producer message goes to 1 partition, and whichever broker is the leader of this partition.

Partition connects to that broker.

Consumer on the other hand can connect to leader or follower node of that partition.



Retention → A key feature of Kafka is that retention which makes it a durable storage of messages for some period of time

topic level retention → (7 days) Remove msgs older than 7 days

Size level retention → (1GB). Partition grows bigger than 1GB, keep partition removing oldest messages

Another new kind of retention policy I learnt about.

(log compacted) \rightarrow Kafka only keeps / retains the last message produced with a specific key

(Example)

offset	key	value	key (101) \rightarrow Partition 1	key (102) \rightarrow Partition 2
0	101	v1		
1	102	v2		
2	101	v3		
3	102	v4		
4	101	v5		

(log compacted)

partition	offset	key	value
1	4	101	v5
2	3	102	v4

This type of retention comes handy in use-cases like

\rightarrow Event Sourcing $\Leftarrow \{$ latest state of cart, profile $\}$

\rightarrow CDC $\Leftarrow \{$ last change log which arrived for a user $\}$
for conflict resolution.

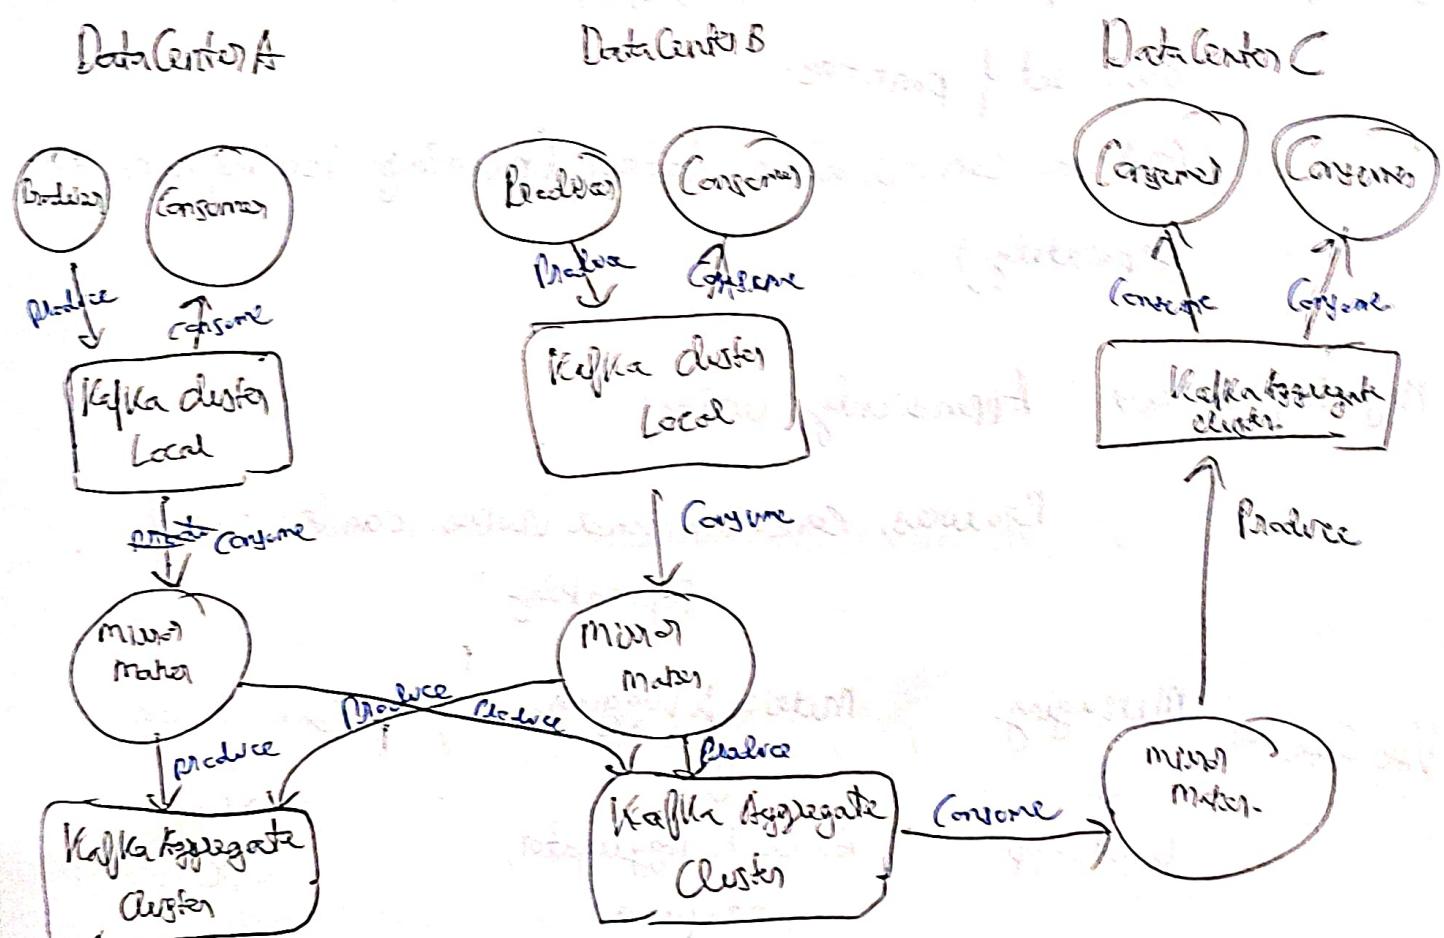
(Multiple Clusters)

The replication mechanisms work within the Kafka cluster and are only designed to work within the Kafka cluster and not b/w multiple clusters.

Kafka project tool, called mirror Maker used for replicating data to other clusters. Mirror Maker is responsible for copying data from one Kafka cluster into another.

At the core Mirror maker is a queue pub/sub only.

One cluster produce to mirror Maker and another cluster consumes from it



(Features of Kafka)

Multiple Producers \rightarrow No need for multiple queries

Multiple Consumers \rightarrow Each consumer has its own offsets, unlike other query

messages consumed by 1 consumer, day it's not make
it available for other consumer.

Other consumer also gets that message from a different partition,
having its own offsets.

Disk-based Retention \rightarrow messages are retained in the Kafka Broker,
and no fear of loss of messages.

Scalable \rightarrow multiple brokers on a cluster, Each broker has its
own set of partitions.

And we can scale a broker (meaning scaling a partition
separately).

High Performance \rightarrow Append only writes,

Producers, Consumers and Brokers can be scaled
separately

Use Cases \rightarrow Messaging, Metrics & Logging,
What'sapp, Ad Click Aggregator,
Grafana

Commit Log
CDC

Birth of Kafka → Kafka developed by LinkedIn's principal engineer who previously developed Voldemort → Distributed K-V Store

Then Kafka → Kafka scaled so high that it was responsible for streaming 5PB data/daily

With horizontal scaling → ↗ Create more partitions, scale consumers, producers, broker separately

Apache Avro for serialisation.

Open Source → 2010 Kafka was made open source,
then it became an Apache Incubator Project
and in 2012 came out as an open source.

Apache Incubator Project → Apache Incubator is the official entry path for new projects to become part of Apache Software Foundation (ASF). From there these projects graduate to become top-level projects.